

What Are the Qualities of Good Malware Datasets?

University of Texas at Dallas (UTD)
Team 2

Faculty Support: Brian Ricks
Sponsor: Frank Posluszny - MITRE

Agenda

- Purpose Statement
- Motivation
- Procedures/Plans
- Schedule
- Project Deliverables
- Limitations and Delimitations
- Detailed Findings
- Issues/Challenges
- Conclusions and Recommendations
- Investigator Biographies
- Bibliography
- Acronyms

Purpose Statement

- **Primary Goal:**

Critically assess methodologies for constructing malware datasets to identify qualities for realistic and effective datasets.

- **Key Activities:**

Identify Limitations: Reviewed existing literature to understand challenges preventing dataset generalization.

Extract Recommendations: Gathered insights from malware research to define qualities of robust datasets.

Learn from Other Fields: Examined disciplines like weather prediction to adopt cross-industry best practices (e.g., regular updates, adaptability).

- **Outcome:**

Comprehensive guidelines for creating unbiased, accurate, and representative malware datasets for research and practical applications.

Motivation

1. Establish Best Practices:

- Develop guidelines for **high-quality malware dataset construction**.

2. Improve Dataset Generalization:

- Ensure datasets are **applicable to real-world threats**.

3. Enhance Detection Tools:

- **Boost the effectiveness** of malware detection and analysis.

4. Cross-Disciplinary Insights:

- Leverage recommendations from other fields to **broaden dataset design strategies**.

Procedures/ Plans

- **Literature Review:** We conducted an extensive review of over 40 papers, focusing on the characteristics of malware datasets and the qualities that define good datasets across various fields.
- **Feature Extraction:** From the reviewed papers, we extracted the strengths, weaknesses, limitations, and recommended features of both malware datasets and datasets from related fields.
- **Cross-Industry Comparison:** We analyzed the dataset recommendations from fields like weather prediction, image processing, and IoT, drawing parallels to the cybersecurity domain, to identify transferable qualities that could enhance malware datasets.

Schedule

Date	Phase	Activity
Aug 27	Project Bids	Submit project bids.
Aug 30	Project Assignments	Receive project assignments.
Sep 8	Project Proposal	Submit project proposal.
Sep 9–15	Phase 1	Research existing malware datasets and limitations.
Sep 16–30	Phase 2	Review cybersecurity papers; identify strengths, limitations, and recommendations.
Oct 1–14	Phase 3	Extract key dataset characteristics (e.g., temporal handling, diversity, representativeness).
Oct 8–14	Phase 4	Analyse cross-disciplinary datasets for applicable practices.
Oct 15–21	Phase 5	Identify best practices for malware dataset creation.
Oct 25–27	Midterm	Deliver midterm presentations and submit deliverables.
Oct 28–Nov 10	Phase 6	Synthesize findings into best practices.
Nov 11–27	Phase 7	Refine best practices; address feedback and gaps.
Nov 28	Final Preparation	Finalize report and presentation.

Project Deliverables

We have three deliverables:

1. **Spreadsheet:** A Google spreadsheet compiling the literature review, feature extraction, and cross-industry analysis. It includes findings from over 40 research papers, key characteristics for effective datasets, and techniques from other fields.
2. **Report:** A detailed document summarizing the findings, methodology, insights, and recommendations for improving malware datasets, based on the literature review and comparative analysis.
3. **Presentation:** Slides highlighting the project's key results, proposed best practices, and recommendations for future malware dataset development.

Limitations and Delimitations

Limitations:

- Reliant on publicly accessible datasets; proprietary data excluded.
- Theoretical focus on methodologies without direct empirical testing.

Delimitations:

- Emphasis on dataset design, not detection algorithms or broader cybersecurity fields.
- Cross-disciplinary comparisons limited to fields like IoT and weather prediction.

Cybersecurity

Key Findings	Description
Historical Coherence	Training data must precede test data to avoid unrealistic performance.
Random Selection Bias	Careful sample selection prevents future data from skewing results.
Model Degradation	Older training data may fail to detect evolving malware threats.
Realistic Construction	Datasets should mimic real-world timelines and conditions of malware appearance.
Static vs. Dynamic Analysis	Static is fast but struggles with obfuscation; dynamic captures runtime behavior but is resource-intensive.

Compiler and Architecture

Key Findings	Description
Environment Dependence	Tools optimized for specific OS may struggle on others (e.g., Windows tools on Linux).
Hardware Dependencies	Tools tailored for specific architectures (e.g., Intel vs. AMD) may limit generalization.
Compiler Variability	Different compilers produce varied binary outputs, impacting analysis.
Compiler Flags	Optimization flags influence binary characteristics and may introduce vulnerabilities.
Cross-Architecture Challenges	Malware may behave differently when recompiled for other architectures.

GraphML and Image Processing

Key Findings	GraphML	Image Processing
Scalability Challenges	Real-world graphs are larger; OGB provides datasets up to 100M nodes.	Datasets should represent realistic conditions.
Lack of Diversity	Structural attributes like node degrees and clustering coefficients are limited in most datasets.	Datasets must avoid biases and include diverse malware samples.
Evaluation Protocols	Standardized evaluation (fixed data splits, metrics) ensures consistent results.	Cross-dataset generalization is key—models must perform well across datasets.
Data Splitting Methods	OGB uses time-based or structure-based splits for more realistic scenarios.	Incorporating realistic negative samples ensures model accuracy and reduces overconfidence.
Handling Rare Events	Include rare but significant graph events.	Dataset should capture rare malware attacks, like APTs and zero-day exploits.

Energy Consumption

Key Findings	Description
Feature Selection	Gradient Boosting identifies impactful features for better model performance in both domains.
Temporal Data	Energy datasets analyze time intervals; malware datasets track attack phases and behaviors.
External Influences	Energy consumption depends on weather; malware detection relies on network and user behavior.
Dataset Diversity	Variations in region, device, and data sources limit generalization in both fields.
Behavioral Insights	Including behavioral data enhances anomaly detection and pattern recognition.
Model Performance	Rich, diverse datasets improve predictions; limited samples hinder generalization.

IoT and Weather Prediction

Key Findings	IoT and Networks	Atmospheric Sciences and Weather Prediction
Dataset Balance	Balancing legitimate and malicious traffic is critical for accurate intrusion detection.	Datasets must capture diverse weather patterns, including rare extreme events like hurricanes.
Feature Selection	Features like message types, QoS levels, and session management are key to detecting subtle attacks.	Selecting critical weather variables (e.g., temperature, wind speed) enhances model accuracy.
Rare Event Inclusion	Attack-specific features improve detection for rare events like APTs.	Include rare but significant weather events to improve predictive models for extreme scenarios.
Model Adaptability	Ensemble models (e.g., Random Forests, RNNs) showed good performance in intrusion detection.	Models need frequent updates to handle evolving climatic conditions and new data sources (e.g., satellite data).

Hardware

Key Findings	Description
Feature Relevance	Focus on microarchitectural features like cache misses and branch mispredictions.
Temporal Context	Include time-based samples to track behavioral changes and detect zero-day threats.
Anomaly Detection	Support signature and anomaly-based methods with sufficient benign examples for modeling.
Diverse Data	Capture varied system behaviors using HPCs across different malware types and benign states.

Issues / Challenges

Lack of Standardized Benchmarks

- The absence of standardized benchmarks for evaluating malware datasets hinders consistent comparison and assessment across different datasets.

Challenges in Adopting Cross-Disciplinary Best Practices

- Adapting best practices from other fields is difficult due to structural differences in data, limiting the transferability of methodologies.

Scalability Limitations

- Techniques effective for small datasets may not perform well with larger datasets, posing scalability challenges in empirical analysis.

Generalizability of Characteristics

- Determining which dataset characteristics can be generalized across different dataset types and use cases remains a complex task.

Conclusions and Recommendations

1. Diversity and Representativeness

- **Malware Type Coverage:** Include ransomware, trojans, worms, APTs, and zero-day exploits.
- **Environmental Representation:** Cover corporate, cloud, consumer, IoT, multiple OS, and architectures.
- **Rare Events:** Include impactful, low-frequency malware (e.g., advanced obfuscation).
- **Attack Strategies:** Capture phishing, SQL injection, DDoS, and credential stuffing.

2. Temporal Coherence and Realism

- **Chronological Splits:** Training data predates testing to simulate real-world scenarios.
- **Lifecycle Modeling:** Track malware behavior from infiltration to execution.
- **Avoid Data Leakage:** Use time-based splits for realistic testing environments.

Conclusions and Recommendations

3. Data Quality and Integrity

- **Precise Annotations:** Accurate labels to reduce false positives/negatives.
- **Noise Reduction:** Remove redundant or misleading attributes.

4. Adaptability and Scalability

- **Frequent Updates:** Incorporate emerging threats with regular updates.
- **Dynamic Corpus:** Use honeypots and crowdsourcing for real-time enrichment.
- **Scalable Design:** Maintain quality as datasets grow.

5. Realism and Practical Application

- **Simulated Real-World Conditions:** Introduce noise, edge cases, and class imbalance.
- **Behavioral Validation:** Include execution traces and system call data.

Conclusions and Recommendations

6. Standardized Benchmarks

- **Evaluation Protocols:** Clear metrics (e.g., accuracy, recall, FPR).
- **Reproducibility:** Publicly available datasets and fixed data splits.
- **Custom Metrics:** Assess resilience to changes (concept drift).

7. Cross-Industry Insights

- **IoT & Network Security:** Emphasize traffic balance and ensemble evaluation.
- **Weather Prediction:** Model rare, impactful events with time-series data.
- **Image Processing:** Reduce biases and include realistic negative samples.

Conclusion

- Emphasize diversity, temporal coherence, and data quality.
- Ensure adaptability, realism, and standardization.
- Leverage cross-industry practices for innovative and robust dataset design.

Biography

- **Mohammad Rafieian:** Ph.D. candidate focused on **advancing software security** through **machine learning-driven vulnerability detection**.
- **Tharun Mulka:** Cybersecurity Ph.D. student specializing in **cryptography** and **innovative defenses against USB-based attacks**.
- **Sahithi Reddy Mopuru:** Master's student excelling in **data science and AI**, with a track record in **healthcare-focused machine learning research**.
- **Sri Kruthi Madugula:** Data Science specialist with strengths in **big data analytics** and **developing predictive models for cybersecurity**.
- **Pranathi Munnangi:** Master's student integrating **AI, cloud computing**, and **cybersecurity**, dedicated to **enhancing safety in intelligent systems**.

Bibliography

1. Kan, Zeliang, et al. "TESSERACT: Eliminating Experimental Bias in Malware Classification across Space and Time (Extended Version)." *arXiv preprint arXiv:2402.01359* (2024).
2. Thompson, Warren, et al. "Malware Dataset Availability & Inherent Bias Study." *2024 IEEE International Conference on Electro Information Technology (eIT)*. IEEE, 2024.
3. Vaccari, Ivan, et al. "MQTTset, a new dataset for machine learning techniques on MQTT." *Sensors* 20.22 (2020): 6578.
4. Dueben, Peter D., et al. "Challenges and benchmark datasets for machine learning in the atmospheric sciences: Definition, status, and outlook." *Artificial Intelligence for the Earth Systems* 1.3 (2022): e210002.
5. Hu, Weihua, et al. "Open graph benchmark: Datasets for machine learning on graphs." *Advances in neural information processing systems* 33 (2020): 22118-22133.
6. Allix, Kevin, et al. "Are your training datasets yet relevant? an investigation into the importance of timeline in machine learning-based malware detection." *International Symposium on Engineering Secure Software and Systems*. Cham: Springer International Publishing, 2015.
7. Pizzolotto, Davide, and Katsuro Inoue. "Identifying compiler and optimization level in binary code from multiple architectures." *IEEE Access* 9 (2021): 163461-163475.
8. Torralba, Antonio, and Alexei A. Efros. "Unbiased look at dataset bias." *CVPR 2011*. IEEE, 2011.
9. Alshaibi, Ahmed, et al. "The comparison of cybersecurity datasets." *Data* 7.2 (2022): 22.

Acronyms

IoT: Internet of Things

APT: Advanced Persistent Threat

OGB: Open Graph Benchmark

RNN: Recurrent Neural Network

SMOTE: Synthetic Minority Over-sampling Technique

PCA: Principal Component Analysis

HPC: Hardware Performance Counters

Terms / Definitions

Malware: Malicious software designed to disrupt, damage, or gain unauthorized access to computer systems.

Dataset Bias: The presence of systematic errors or skewness in data collection, leading to a dataset that does not fully represent the diversity or complexity of real-world malware.

Dataset Quality: The criteria by which a dataset is evaluated, including factors such as accuracy, completeness, consistency, and relevance.

Real-world Malware: Malware samples that are observed in actual cyber-attacks or infections, as opposed to artificial or laboratory-created samples.

Data Diversity: The range and variety of malware samples included in a dataset to ensure a comprehensive representation of different malware types and behaviors.

Malware Detection: The process of identifying malware based on its signatures, behaviors, or characteristics within a system or dataset.

False Positives/Negatives: Incorrect results in malware detection, where benign data is identified as malware (false positive) or malware goes undetected (false negative).

A representative malware dataset is a diverse and comprehensive collection of malware samples that accurately reflects the variety and complexity of real-world threats, ensuring effective analysis and detection across different malware types.



THANK YOU