

INSuRE Fall 2024

Midterm Document - October, 2024

What Are the Qualities of “Good” Malware Datasets?

Investigators:

The University of Texas at Dallas

Team 2

Sai Tharun Reddy Mulka - Student researcher - sxm220351@utdallas.edu

Mohammad Rafeian - Student Researcher - mxr230010@utdallas.edu

Pranathi Munnangi - Student - pxm220084@utdallas.edu

Sahithi Reddy Mopuru - Student - sxm220353@utdallas.edu

Sri Kruthi Madugula - Student - sxm220429@utdallas.edu

Brian Ricks - Faculty Support - absolutefunk@utdallas.edu

Problem Mentor:

Frank Posluszny - MITRE - malwarefrank@mitre.org

1. Project Introduction	2
1.1. Research Project Description	2
1.2. Project Objectives / Goals	3
1.3. Investigator Biographies	3
2. Research Plan	6
2.1. Research Plan Summary	6
Research objective	6
Division of tasks	7
Methodology	7
Expected outcome	7
Schedule	8
2.2. Deliverables	8
3. Progress To Date	8
Cyber	8
Image Processing and GraphML	9
Building and Compilers	11
IOT and Networks	12
Atmospheric Sciences and Weather prediction	13
4. Summary of Results / Conclusions to Date	15
5. Issues and Challenges	16
6. Bibliography	16
Appendix A: Acronyms and Terms / Definitions	16

1. Project Introduction

1.1. Research Project Description

This project aims to establish the key characteristics of a "good" malware dataset by thoroughly examining the limitations and biases present in existing datasets. Many current datasets are biased due to the ease of collecting certain types of data, which leads to an inaccurate representation of real-world malware behaviors and characteristics. This skewed data can significantly affect the validity of research findings, making it difficult to generalize results beyond controlled environments.

To address these challenges, the project will conduct an in-depth evaluation of current practices in malware dataset construction, drawing from research in both cybersecurity and other fields where data realism and diversity are critical. By leveraging cross-disciplinary insights, the project seeks to develop a set of best practices and guidelines for building malware datasets that accurately capture the range of features and behaviors exhibited by real-world threats. The goal is to create datasets that are comprehensive, diverse, and applicable across various research and practical scenarios, thereby enhancing the quality and reliability of future malware studies.

1.2. Project Objectives / Goals

The primary goal of this project is to critically evaluate the current methodologies used to construct malware datasets and to identify the characteristics that make a dataset both realistic and effective for research and practical applications. This involves a deep dive into existing literature and practices in malware research as well as other disciplines that deal with complex, real-world data. For more advanced student groups, the project may also include conducting a survey among malware analysis researchers and practitioners to gather first-hand insights on the qualities they believe are essential in a "good" malware dataset. The ultimate aim is to produce a comprehensive set of guidelines and recommendations that can be used by the research community to develop more accurate, unbiased, and representative malware datasets.

1.3. Investigator Biographies

Mohammad Rafieian is a Ph.D. candidate in Software Engineering at The University of Texas at Dallas, with a focus on software security and machine learning. His research aims to enhance the efficiency of automated program analyzers in detecting vulnerabilities by applying advanced machine learning methods. As a Research Assistant, Mohammad has contributed to the development and testing of innovative program analysis tools, playing a pivotal

role in advancing cybersecurity. In addition to his research, he has worked as a Backend Developer, where he designed scalable architectures and optimized system performance. Proficient in Python, Java, and Docker, Mohammad's combination of cutting-edge security research and hands-on development experience makes him a valuable contributor to any cybersecurity Project.

Tharun Mulka is a second-year Ph.D. student at The University of Texas at Dallas, specializing in cybersecurity, with a strong foundation in both academic research and practical applications. He holds a bachelor's degree in computer science, where his research on email security led to the discovery of vulnerabilities in the DKIM authentication protocol, sharpening his problem-solving and analytical skills. His experience publishing research has strengthened his academic writing, while his advanced coursework in Network Security, Algorithms, and Information Security has deepened his technical knowledge. Tharun has demonstrated his ability to build complex systems from scratch, such as implementing end-to-end encryption for an e-commerce platform without relying on external libraries, showcasing his expertise in cryptography. His current Ph.D. research focuses on securing data and privacy against USB-based attacks, which has given him hands-on experience with the USB protocol and further enhanced his coding skills. Given his diverse experience and ability to quickly master new concepts, Tharun is well-equipped to excel in his project and contribute effectively to the team's success.

Sahithi Reddy Mopuru is currently pursuing her Master of Science in Computer Science at The University of Texas at Dallas, building on a strong foundation established during her undergraduate studies in computer science and engineering at SRM University, Chennai. With an impressive academic background, Sahithi has participated in several significant research projects that have honed her technical and analytical skills. During her bachelor's, she developed a keen interest in data science and machine learning, which culminated in a project focused on detecting counterfeit medical remedies using an LSTM model in the TensorFlow framework. This work, presented at the IEEE 12th International Conference on Data Science and Engineering, gave her valuable insights into healthcare misinformation and highlighted her commitment to leveraging technology for societal impact. In her graduate studies, Sahithi has expanded her expertise in data science, excelling in data analysis, machine learning, and statistical modeling with tools like Python, TensorFlow, and R. Her technical skills were further demonstrated through her development of a smart fish farming recommendation system, which used reinforcement learning algorithms to optimize feeding schedules and enhance fish health. In addition to her work in data science, Sahithi has cultivated a strong interest in cybersecurity, focusing on understanding and mitigating digital threats. This combination of data science and cybersecurity skills equips her to solve complex,

real-world problems with innovative solutions. Currently, as a Grading Assistant at UT Dallas, Sahithi continues to refine her technical abilities by providing data analysis support and personalized feedback to students, contributing to both her research pursuits and professional growth.

Sri Kruthi Madugula is a Master of Science student in Computer Science with a specialization in Data Science at The University of Texas at Dallas. Her academic background, combined with hands-on experience, makes her a valuable contributor to cybersecurity projects. Sri Kruthi has developed a strong foundation in data structures, algorithms, machine learning, and big data management—skills essential for analyzing and mitigating cyber threats. Her proficiency in designing and implementing complex algorithms is especially relevant for building robust malware detection systems that require advanced data processing and analysis to identify potential security breaches. During her tenure as a Data Engineering Intern at Tenet HealthCare, Sri Kruthi gained significant experience managing and securing large-scale data environments. She successfully migrated SQL Server stored procedures to Google BigQuery, demonstrating her expertise in handling sensitive data, ensuring security, and optimizing systems for real-time analytics—skills crucial for maintaining data integrity in cybersecurity. Her work on machine learning models for liver disease prediction and grape disease detection further highlights her ability to apply advanced analytics to solve complex problems. In the cybersecurity context, these abilities are transferable to developing predictive models for threat detection and response. With a combination of strong academic knowledge and practical engineering experience, Sri Kruthi brings a unique blend of data science expertise to any cybersecurity-focused team. Her capabilities in data-driven approaches make her a key player in projects aimed at enhancing digital security.

Pranathi Munnangi is a second-year Master's student in Computer Science, specializing in Data Science. She earned her Bachelor's degree in Computer Science, with a focus on cloud computing, cybersecurity, network security, and AI. This strong foundation has provided her with a deep understanding of securing systems, protecting sensitive data, and deploying secure cloud infrastructures. In her current studies, Pranathi has developed expertise in advanced data science techniques, including machine learning, data mining, and big data technologies. She is proficient in programming languages such as Python, R, C++, and Java, and has experience working with SQL and Spark for data processing and analysis. Her interest in cloud computing has led her to explore the integration of AI and machine learning in cloud environments, with a focus on optimizing performance and enhancing security. In her project, "Generating Images Using Complex Sentences," Pranathi investigated the capabilities of Stable Diffusion models, a cutting-edge technology in generative AI. This diverse background in AI, cloud computing, and cybersecurity, combined with her practical programming and networking experience, equips Pranathi to

address challenges at the intersection of these fields. Pranathi is dedicated to integrating robust safety protocols into AI models and cloud systems, ensuring that these technologies operate reliably and ethically. Her ultimate goal is to develop AI and cloud computing tools that enhance human capabilities while minimizing potential risks, reflecting her commitment to creating a more secure and trustworthy technological future.

2. Research Plan

In our research on identifying the qualities of good malware datasets, we evaluated over 30 papers, extracting both the limitations and recommendations for existing datasets. Malware datasets are often designed for specific projects, which limits their generalizability. Each dataset in the cybersecurity domain has its own strengths and weaknesses. We also explored other fields, which offer valuable insights on dataset qualities that, if effectively generalized, could benefit malware dataset design.

Identifying Limitations of Malware Datasets

Our first task was to review the limitations highlighted in malware dataset papers. This step was crucial for identifying the characteristics that make these datasets less applicable across different projects, helping us understand what prevents them from being generalized.

Extracting Recommendations for Malware Datasets

Next, through a careful literature review, we gathered recommendations on the qualities that define good malware datasets. These recommendations provide insights into how future datasets can be improved for broader use in cybersecurity.

Expanding Insights from Other Fields

In addition to cybersecurity, we examined other fields like weather prediction to extract recommendations for dataset qualities. These fields share characteristics with malware datasets, such as the need for regular updates and unpredictability of future features. We argue that by incorporating recommended qualities from these fields, we can design more robust and versatile malware datasets.

This approach of cross-industry evaluation provides a more holistic view of how to overcome existing limitations and improve dataset generalization in malware research.

2.1. Research Plan Summary

Research objective

The primary objective of our research is to identify the key qualities, recommendations, and limitations of existing malware datasets, and to propose a comprehensive set of best practices for creating future malware datasets. By conducting an in-depth literature review of 30 papers from both the cybersecurity field and related disciplines, we aim to extract the essential features that define effective datasets. We argue that insights from other fields, such as weather prediction and image processing, can provide valuable guidance in addressing the challenges of malware dataset design and improving their generalizability.

Division of tasks

Our research is structured into three key tasks:

1. **Literature Review:** We conducted an extensive review of over 30 papers, focusing on the characteristics of malware datasets and the qualities that define good datasets across various fields.
2. **Feature Extraction:** From the reviewed papers, we extracted the strengths, weaknesses, limitations, and recommended features of both malware datasets and datasets from related fields.
3. **Cross-Industry Comparison:** We analyzed the dataset recommendations from fields like weather prediction, image processing, and IoT, drawing parallels to the cybersecurity domain, to identify transferable qualities that could enhance malware datasets.

Methodology

Our research follows a qualitative methodology that includes:

- **Systematic Literature Review:** We examined a wide range of research papers on malware datasets, compiling data on limitations and recommendations.
- **Comparative Analysis:** We evaluated how other fields, which deal with similar challenges (e.g., unpredictability, constant updating), structure and generate their datasets.
- **Synthesis of Findings:** The final step involves synthesizing insights from both cybersecurity and other fields to propose a set of best practices and recommendations for future malware datasets.

Expected outcome

The expected outcome of our research is the formulation of a set of comprehensive qualities and recommendations for future malware dataset creation. We anticipate that our findings will highlight common dataset shortcomings and suggest ways to address them. Moreover, we expect that features and dataset recommendations from other fields—such as the importance of regular updates and handling of diverse environments—will provide valuable strategies to improve the robustness, scalability, and generalizability of future malware datasets.

Schedule

Phase	Date	Programming Group
	August 27	Project Bids
	August 30	Project Assignments
	September 8	Project Proposal
Phase 1	September 9 - September 15	Initial research on existing malware datasets and their limitations
Phase 2	September 16 - September 31	Systematically review research papers from cybersecurity and related fields. Identify and categorize the limitations, strengths, and recommendations for malware datasets.
Phase 3	October 1 - October 14	Extract key characteristics that define good malware datasets from the reviewed papers. Focus on features like temporal data handling, dataset diversity, and representativeness.
Phase 4	October 8 - October 14	Analyze cross-disciplinary datasets, identifying applicable techniques and qualities from fields like atmospheric sciences and image processing.
Phase 5	October 15 - October 21	finding good practices for

		creating datasets/ malware datasets
	October 25 - October 27	Midterm presentations / Midterm deliverables
Phase 6	October 28 - November 10	Synthesize findings from the literature review and comparative analysis to develop best practices for malware dataset creation.
Phase 7	November 11 - November 27	Review and refine the best practices based on feedback. Address any gaps or additional considerations before finalizing deliverables.
	November 28	Final report writing and presentation practice run

2.2. Deliverables

We have three deliverables:

1. **Spreadsheet:** A Google spreadsheet compiling the literature review, feature extraction, and cross-industry analysis. It includes findings from over 30 research papers, key characteristics for effective datasets, and techniques from other fields.
2. **Report:** A detailed document summarizing the findings, methodology, insights, and recommendations for improving malware datasets, based on the literature review and comparative analysis.
3. **Presentation:** Slides highlighting the project's key results, proposed best practices, and recommendations for future malware dataset development.

3. Progress To Date

We divided our findings into different categories according to the recommendations of different industries.

Cyber

Characteristic	Description	Limitations/Considerations
Historical Coherence	Training data should precede test data chronologically to avoid bias.	Using future data for training can lead to inflated performance results.
Random Selection Bias	Avoid randomly selecting training data without considering the timeline.	May incorporate "future" malware in training, skewing results.
Model Degradation Over Time	Detection model performance decreases as the gap between training and test datasets increases.	Older training data may not represent new malware threats.
Diverse Datasets	Training on a broad set of malware families improves generalization.	Limited diversity in training data can result in poor performance on unseen or rare threats.
Realistic Construction	Datasets should mimic real-world timelines and conditions of malware appearance.	Unrealistic datasets can lead to poor real-world performance.
Measure Multiple Performance Metrics	Evaluate metrics beyond runtime, such as detection accuracy, false positives/negatives, and memory usage.	Focusing solely on runtime can neglect important aspects of model effectiveness.
Raw and Processed Features	Raw features allow custom feature engineering, while processed features help focus on algorithm optimization.	Using only one type may limit experimentation or model performance.
Class Imbalance	Malware datasets often have a higher number of benign samples than malicious ones.	Can skew evaluation metrics if not properly addressed (e.g., using SMOTE or balancing techniques).
Realistic Validation ("In the Wild")	Real-world testing ensures the model's effectiveness outside the lab environment.	Lab settings may not cover all scenarios encountered in the wild.

Select Informative Features	Focus on features that add value to detection while avoiding overfitting and high computational costs.	Including too many features can lead to overfitting or slow performance.
Handling Adware	Adware detection is challenging due to its ambiguous nature, lying between benign and malicious categories.	Can affect model accuracy; careful labeling and feature selection are needed.
Static Analysis	Extracts features from Portable Executable (PE) files without running the code (e.g., section size, entropy).	Struggles with packed/obfuscated malware; techniques like encryption can hinder feature extraction.
Dynamic Analysis	Observes malware behavior in execution (e.g., system calls, network activity).	Limited execution paths may not reveal all behaviors; anti-analysis techniques may evade detection.

Image Processing and GraphML

The **GraphML** the key findings are :

1. **Scalability Challenges:** Commonly used graph-based datasets tend to be much smaller than the graphs encountered in real-world applications, limiting the evaluation of models for large-scale scenarios. There is a need for datasets that are sufficiently large to enable training on complex models and ensure their ability to generalize effectively. The OGB addresses this by offering datasets of various sizes, ranging from smaller graphs to those with up to 100 million nodes, thereby supporting scalability testing across different scales.
2. **Lack of Diversity:** Many datasets lack variety in structural attributes such as node degrees, clustering coefficients, and other graph characteristics, which restricts a model's ability to generalize across different domains. The OGB datasets strive to encompass graphs from a broad spectrum of domains, including social and information networks, biological networks, molecular structures, source code abstract syntax trees (ASTs), and knowledge graphs. This extensive range of domains supports the creation and evaluation of versatile models, allowing for a clear distinction between general-purpose approaches and domain-specific methods.
3. **Inconsistent Evaluation Protocols:** The absence of standardized evaluation protocols across studies has led to inconsistent results, making comparisons challenging. Using fixed data splits and unified evaluation metrics can standardize the evaluation process, improve reproducibility, and allow for more meaningful comparisons of different approaches.

4. **Data Splitting Methods:** Random dataset splitting is commonly used but does not accurately reflect real-world scenarios, where time-based or structurally informed splits are more applicable. OGB employs more realistic splitting methods (e.g., time-based or molecular structure-based splits) to better simulate real-world conditions, providing more practical and meaningful evaluation scenarios.
5. **Balanced Representation:** Balanced datasets are crucial for preventing bias in model training and achieving fair evaluations. Some OGB datasets, particularly those used in molecular studies, are curated to address class imbalance challenges, and evaluation metrics like Average Precision (AP) are employed to account for skewed data distributions.

These findings underscore the necessity for enhanced benchmarks that reflect real-world applications, facilitate reproducible evaluations, and provide consistent frameworks for machine learning research in the context of graph data.

In **Image processing** the key findings are :

The key findings from the paper on good characteristics for malware datasets include:

1. **Realism and Practical Relevance:** Datasets should represent realistic scenarios to ensure practical applicability in real-world situations. Simulated environments or overly controlled datasets may not adequately prepare models for real-life conditions.
2. **Bias Mitigation:** It's essential to identify and minimize biases in the dataset, such as selection, capture, and negative set biases. Ensuring diverse sampling from various sources can help reduce these biases.
3. **Cross-Dataset Generalization:** A high-quality dataset should enable generalization across different datasets, showing minimal performance drop when tested in different environments. This characteristic helps validate the dataset's representativeness of the broader visual world.
4. **Negative Set Inclusion:** Incorporating diverse and realistic negative samples (instances where malware is not present) is crucial for classifier training, ensuring models do not become overconfident and can better distinguish between positive and negative cases.
5. **Training Data Sufficiency:** The dataset should include enough training examples to allow models to learn effectively. Ideally, the dataset size should be large enough to avoid overfitting while capturing the diversity in malware characteristics.
6. **Annotation Quality:** The quality and consistency of labels (annotations) in the dataset significantly affect model performance. Consistent, high-quality annotations ensure accurate training and evaluation.
7. **Data Augmentation and Transformation:** To address capture biases, applying data transformations such as flipping, jittering, or cropping can help create a more robust training dataset.

Building and Compilers

In the field of malware analysis, a notable limitation is that most existing datasets are tailored for specific environments and operating systems, primarily focusing on Windows. This specificity hampers reproducibility across diverse platforms, highlighting the need for comprehensive datasets that encompass various environments, including Windows, Linux, and macOS. Additionally, hardware specifications play a crucial role in the build process and data generation, particularly when utilizing tools like Intel Pin, which analyzes binaries reliably only on Intel CPU-based systems. While Intel Pin is effective for analyzing Windows malware, its reliance on Intel CPUs necessitates re-implementing analysis tools like EvDetector to accommodate other CPU architectures, ensuring broader applicability and reproducibility of findings.

Moreover, the generation of binary files is influenced by different compilers and optimization flags; however, these critical details are not embedded within the binaries themselves. Understanding compiler and flag configurations is essential for identifying vulnerabilities, rewriting binary code, and optimizing energy consumption. Recent studies have proposed datasets that include binary files alongside their associated compiler and optimization flags. These datasets are employed to train neural networks for predicting optimization levels across various architectures. The authors emphasize the importance of including features such as architecture, compiler, and compiler flags in dataset design, enabling more effective training and improved model accuracy.

Feature	Importance	Application
Environment	Datasets often focus on specific OS (e.g., Windows), limiting reproducibility.	Need for datasets covering various OS (Windows, Linux, macOS) for broader applicability.
Hardware Specs	Intel Pin analyzes binaries reliably on Intel CPUs, limiting analysis to Intel architecture.	Re-implementation of tools (e.g., EvDetector) is required for compatibility with other CPUs.
Compiler	Different compilers produce varied binary outputs.	Essential for vulnerability discovery and performance analysis.
Compiler Flags	Flags can determine performance and consumption.	Important for optimizing binary rewriting and improving runtime efficiency.

Architecture Affects how binaries are executed and analyzed. Critical for training models that predict optimization levels across different architectures.

IOT and Networks

In the context of IoT and network security, these key findings highlight how the dataset contributes to improving accuracy in machine learning models for intrusion detection and what kind of features are critical, especially for attack detection.

1. Dataset Design and Balance for Accuracy

- Issue with Traffic Dominance: A critical finding is that the dataset includes an overwhelmingly large volume of legitimate traffic compared to malicious traffic. This imbalance can skew accuracy metrics, as models may primarily learn from the abundant legitimate traffic, leading to potential misclassification of rare attack types.
- Balancing legitimate and malicious traffic is crucial for ensuring that the models do not overfit to benign patterns and can generalize to detect rare attack instances effectively. Datasets like BoT-IoT and TON_IoT show that scalable datasets with long-term traffic monitoring improve the detection of low-traffic, high-persistence attacks (e.g., data exfiltration, botnet communications).

2. Feature Selection and Importance

- In IoT datasets like MQTTset, N-BaloT and IoT-23 highlight flow-level features and DNS analysis for botnet detection, MQTTset focuses on timing and message types for detecting DoS and SlowITe attacks. The features extracted focus heavily on the characteristics of the MQTT protocol, including both network-layer features (TCP/IP) and application-layer MQTT features. These include flags, session information, message types, QoS levels, Timing-Based, Traffic Length, Size and other protocol-specific data.

3. Attack-Specific Features and Detection

- Each type of attack has specific protocol-related features that are critical for detection. Flooding attacks rely on traffic volume and message repetition features, while more subtle attacks like SlowITe depend on timing and session management features. While designing a dataset targeting attack specific features will help the model to detect well.

4. Machine Learning Model Performance

- Ensembled based approach gave good results, It worked well in detecting anomalies by analyzing multi-dimensional features like message types, QoS levels, and timing differences. BoT-IoT employed traditional machine learning models like random forests

and gradient boosting supports the usage of ensemble models. For time-series anomaly detection TON_IoT showed RNN, Autoencoders worked very well.

Atmospheric Sciences and Weather prediction

1. Data Diversity and Representativeness

In both atmospheric sciences and weather prediction, datasets need to capture a wide range of environmental conditions, geographic regions, and time periods to ensure models can generalize across different weather patterns, including extreme events. This diversity allows for more accurate and robust predictions. Similarly, a good malware dataset** should encompass a variety of malware types, attack strategies, operating systems, and real-world scenarios. This ensures the model can generalize effectively across different cyber threats, making it adaptable and reliable in detecting diverse and evolving malware threats. The emphasis on diversity and representativeness in both fields helps create models that are better prepared for real-world challenges.

2. Temporal Data and Sequence Modeling

In both atmospheric sciences and weather prediction, time-series data is essential for modeling how conditions evolve over time, whether hourly, daily, or seasonally. This sequential data helps models accurately predict future weather patterns by understanding changes in variables like temperature, pressure, and wind speed. Similarly, a good malware dataset should capture the temporal sequence of system events, network activity, and user behavior to model the progression of an attack. By tracking how malware behaviors unfold over time—from initial infiltration to exploitation—these datasets enable real-time threat detection and help models adapt to evolving attack patterns.

3. Handling Rare and Extreme Events

In both atmospheric sciences and weather prediction, it is crucial to include rare but significant events, such as hurricanes, tornadoes, or flash floods, in datasets. These extreme occurrences can greatly influence predictions and outcomes, so capturing them ensures models are not just trained on everyday weather patterns but can also accurately predict and respond to high-impact situations. Similarly, a **good malware dataset** should encompass rare but critical attack types, such as advanced persistent threats (APTs) or zero-day exploits. Although these attacks are infrequent, their potential impact is substantial. Including such rare events in the dataset ensures that models are robust and equipped to detect unusual but highly consequential threats, thereby enhancing their resilience and real-world applicability.

4. Feature Selection and Importance

In both atmospheric sciences and weather prediction, selecting key variables like temperature, wind speed, and humidity is essential for accurate predictions, as the choice of features directly influences the quality of the model. Identifying the most predictive atmospheric variables, such as air pressure changes for storm forecasts,

refines the models and improves their performance. Similarly, a good malware dataset requires careful selection of important features, such as API calls, network traffic patterns, and system events. Focusing on the most relevant features enhances the effectiveness of detection models by improving accuracy and reducing noise, making them more capable of identifying threats.

5. Benchmarking and Standardization

In both atmospheric sciences and weather prediction, benchmark datasets and standardized evaluation metrics are essential for consistently assessing model accuracy and reliability across different conditions. Standardized weather datasets, such as those from the European Centre for Medium-Range Weather Forecasts (ECMWF), enable meaningful comparisons of forecasting approaches. Similarly, in malware research, having benchmark datasets allows for fair comparisons of detection techniques under consistent conditions. This standardization helps improve model development by providing a clear basis for evaluating and refining different detection approaches, ensuring that models are robust and effective across various scenarios.

6. Adaptability and Update Frequency

In both atmospheric sciences and weather prediction, regularly updating datasets is crucial to account for evolving climatic conditions and new data sources, such as satellite observations, ensuring models remain accurate and adaptable to changes. Similarly, for a good malware dataset, frequent updates are essential to incorporate new malware types, evolving attack methods, and emerging security threats. This continuous updating process helps maintain the effectiveness of detection models, enabling them to adapt to the rapidly changing cyber threat landscape and ensuring robust, real-time threat detection.

7. Data Quality and Integrity

In both atmospheric sciences and weather prediction, ensuring high data quality is essential for accurate model predictions. This involves reducing noise, correcting errors, and ensuring data completeness through reliable data collection methods, such as satellite observations and ground stations. Similarly, for a good malware dataset, maintaining data quality means filtering out noise, ensuring accurate labeling, and minimizing false positives. High data integrity is crucial in both fields, as it directly affects the reliability and trustworthiness of the models, enabling them to make accurate and effective predictions or detections.

4. Summary of Results / Conclusions to Date

In our literature review, we identified that existing malware datasets are often too specialized and cannot be generalized across different projects. Instead of updating and building upon previous datasets, researchers tend to develop new ones for each specific task, resulting in an abundance of datasets that lack interchangeability. Through a careful evaluation of these works,

we have distilled the most essential features of malware datasets as recommended by researchers. Additionally, we broadened our investigation by exploring other fields that share characteristics with malware analysis, identifying key features that could be beneficial for malware datasets.

For example, malware analysis requires regular updates and anticipation of unseen features—similar to weather prediction, where data must be updated frequently for accurate forecasting by machine learning models. In total, we examined 30 papers across various fields to identify both the limitations of current cybersecurity datasets and recommendations for improvement. These fields include:

- Cybersecurity
- Image Processing and Graph Neural Networks
- Compilers and Program Analysis
- IoT and Network Security
- Weather Prediction

By drawing insights from these diverse fields, we provide a more comprehensive perspective on the essential features needed to enhance the generalization and applicability of malware datasets.

5. Issues and Challenges

- **Lack of Standardized Benchmarks:** The lack of standardized benchmarks makes it difficult to evaluate and compare different malware datasets consistently. Without common evaluation metrics, determining which characteristics truly contribute to dataset quality becomes subjective.
- **Challenges in Adopting Cross-Disciplinary Best Practices:** While insights from fields like weather prediction and image processing are valuable, adapting their best practices to malware datasets presents unique challenges due to fundamental differences in data structure and characteristics.
- **Scalability Limitations:** Methods that work effectively for analyzing smaller datasets often face significant challenges when applied to larger-scale malware collections. This scalability issue affects both the analysis process and the validation of findings.
- **Accounting for Emerging Threats:** The constant evolution of malware threats makes it challenging to ensure that identified characteristics remain relevant and applicable to emerging threats and attack patterns.
- **Generalizability of Characteristics:** Identifying which characteristics of a dataset can be consistently applied across various dataset types and use cases is challenging, as different datasets may have unique properties, biases, or specific features that limit the generalization of findings.

6. Bibliography

- Thompson, W., Elammas, T., Kalappattil, M., Rizi, A., Sunderland, G., Baldizon, L., Posluszny, F. and Wang, Y., 2024, May. Malware Dataset Availability & Inherent Bias Study. In *2024 IEEE International Conference on Electro Information Technology (eIT)* (pp. 202-208). IEEE.
- Kan, Z., McFadden, S., Arp, D., Pendlebury, F., Jordaney, R., Kinder, J., Pierazzi, F. and Cavallaro, L., 2024. TESSERACT: Eliminating Experimental Bias in Malware Classification across Space and Time (Extended Version). *arXiv preprint arXiv:2402.01359*.
- Vaccari, I., Chiola, G., Aiello, M., Mongelli, M. and Cambiaso, E., 2020. MQTTset, a new dataset for machine learning techniques on MQTT. *Sensors*, 20(22), p.6578.
- Le Quy, T., Roy, A., Iosifidis, V., Zhang, W. and Ntoutsis, E., 2022. A survey on datasets for fairness-aware machine learning. *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery*, 12(3), p.e1452.
- Hu, W., Fey, M., Zitnik, M., Dong, Y., Ren, H., Liu, B., Catasta, M. and Leskovec, J., 2020. Open graph benchmark: Datasets for machine learning on graphs. *Advances in neural information processing systems*, 33, pp.22118-22133.
- Dueben, P.D., Schultz, M.G., Chantry, M., Gagne, D.J., Hall, D.M. and McGovern, A., 2022. Challenges and benchmark datasets for machine learning in the atmospheric sciences: Definition, status, and outlook. *Artificial Intelligence for the Earth Systems*, 1(3), p.e210002.
- Alshaibi, A., Al-Ani, M., Al-Azzawi, A., Konev, A. and Shelupanov, A., 2022. The comparison of cybersecurity datasets. *Data*, 7(2), p.22.
- Larriva-Novo, X.A., Vega-Barbas, M., Villagr , V.A. and Rodrigo, M.S., 2020. Evaluation of cybersecurity data set characteristics for their applicability to neural networks algorithms detecting cybersecurity anomalies. *IEEE Access*, 8, pp.9005-9014.
- Allix, K., Bissyand , T.F., Klein, J. and Le Traon, Y., 2015, March. Are your training datasets yet relevant? an investigation into the importance of timeline in machine learning-based malware detection. In *International Symposium on Engineering Secure Software and Systems* (pp. 51-67). Cham: Springer International Publishing.
- Anderson, H.S. and Roth, P., 2018. Ember: an open dataset for training static pe malware machine learning models. *arXiv preprint arXiv:1804.04637*.
- Raman, K., 2012. Selecting features to classify malware. *InfoSec Southwest, 2012*, pp.1-5.
- Ijaz, M., Durad, M.H. and Ismail, M., 2019, January. Static and dynamic malware analysis using machine learning. In *2019 16th International bhurban conference on applied sciences and technology (IBCAST)* (pp. 687-691). IEEE.
- Sommer, R. and Paxson, V., 2010, May. Outside the closed world: On using machine learning for network intrusion detection. In *2010 IEEE symposium on security and privacy* (pp. 305-316). IEEE.
- van der Kouwe, E., Andriesse, D., Bos, H., Giuffrida, C. and Heiser, G., 2018. Benchmarking crimes: an emerging threat in systems security. *arXiv preprint arXiv:1801.02381*.
- Allix, K., Bissyand , T.F., J rome, Q., Klein, J., State, R. and Le Traon, Y., 2016. Empirical assessment of machine learning-based malware detectors for Android: Measuring the gap between in-the-lab and in-the-wild validation scenarios. *Empirical Software Engineering*, 21, pp.183-211.
- Roy, S., DeLoach, J., Li, Y., Herndon, N., Caragea, D., Ou, X., Ranganath, V.P., Li, H. and Guevara, N., 2015, December. Experimental study with real-world data for android app security analysis using machine learning. In *Proceedings of the 31st Annual Computer Security Applications Conference* (pp. 81-90).

- Jordaney, R., Sharad, K., Dash, S.K., Wang, Z., Papini, D., Nouretdinov, I. and Cavallaro, L., 2017. Transcend: Detecting concept drift in malware classification models. In *26th USENIX security symposium (USENIX security 17)* (pp. 625-642).
- Torralba, A. and Efros, A.A., 2011, June. Unbiased look at dataset bias. In *CVPR 2011* (pp. 1521-1528). IEEE.
- Khosla, A., Zhou, T., Malisiewicz, T., Efros, A.A. and Torralba, A., 2012. Undoing the damage of dataset bias. In *Computer Vision–ECCV 2012: 12th European Conference on Computer Vision, Florence, Italy, October 7-13, 2012, Proceedings, Part I 12* (pp. 158-171). Springer Berlin Heidelberg.
- Tommasi, T., Patricia, N., Caputo, B. and Tuytelaars, T., 2017. A deeper look at dataset bias. *Domain adaptation in computer vision applications*, pp.37-55.
- Pizzolotto, D. and Inoue, K., 2021. Identifying compiler and optimization level in binary code from multiple architectures. *IEEE Access*, 9, pp.163461-163475.
- Borah, P., Bhattacharyya, D.K. and Kalita, J.K., 2020, December. Malware dataset generation and evaluation. In *2020 IEEE 4th Conference on Information & Communication Technology (CICT)* (pp. 1-6). IEEE.

Appendix A: Acronyms and Terms / Definitions

PE : Portable Executable File

APT: Advanced Persistent Threat

OGB: Open Graph Benchmark

TCP/IP: Transmission Control Protocol/Internet Protocol

RNN: Recurrent Neural Network

DKIM: Domain Keys Identified Mail

SMOTE: Synthetic Minority Over-sampling Technique

QOS: Quality of Service

MQTT protocol: Message Queuing Telemetry Transport is a lightweight, publish-subscribe protocol designed for efficient communication in IoT and low-bandwidth networks.

Malware: Malicious software designed to disrupt, damage, or gain unauthorized access to computer systems.

Dataset Bias: The presence of systematic errors or skewness in data collection, leading to a dataset that does not fully represent the diversity or complexity of real-world malware.

Dataset Quality: The criteria by which a dataset is evaluated, including factors such as accuracy, completeness, consistency, and relevance.

Real-world Malware: Malware samples that are observed in actual cyber-attacks or infections, as opposed to artificial or laboratory-created samples.

Data Diversity: The range and variety of malware samples included in a dataset to ensure a comprehensive representation of different malware types and behaviors.

Malware Detection: The process of identifying malware based on its signatures, behaviors, or characteristics within a system or dataset.

False Positives/Negatives: Incorrect results in malware detection, where benign data is identified as malware (false positive) or malware goes undetected (false negative).

A **representative malware dataset** is a diverse and comprehensive collection of malware samples that accurately reflects the variety and complexity of real-world threats, ensuring effective analysis and detection across different malware types.